

# Cancer Prediction

AUTHOR

Qinglan Ouyang

```
# 加载包  
library(tidyverse)
```

Warning: package 'forcats' was built under R version 4.4.3

```
— Attaching core tidyverse packages ————— tidyverse 2.0.0 —  
✓ dplyr     1.1.4      ✓ readr     2.1.5  
✓ forcats   1.0.1      ✓ stringr   1.5.1  
✓ ggplot2   3.5.1      ✓ tibble    3.2.1  
✓ lubridate 1.9.3      ✓ tidyr    1.3.1  
✓ purrr    1.0.2  
— Conflicts ————— tidyverse_conflicts() —  
✖ dplyr::filter() masks stats::filter()  
✖ dplyr::lag()   masks stats::lag()  
ℹ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become  
errors
```

```
library(caret)      # 用于模型训练和交叉验证
```

Warning: package 'caret' was built under R version 4.4.3

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

```
library(e1071)      # 用于SVM和Naive Bayes
```

Warning: package 'e1071' was built under R version 4.4.3

```
library(pROC)        # 用于ROC曲线和AUC
```

Warning: package 'pROC' was built under R version 4.4.3

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:stats':

```
cov, smooth, var
```

```
library(doParallel) # 用于并行处理 (可选, 加速CV)
```

Warning: package 'doParallel' was built under R version 4.4.3

Loading required package: foreach

Warning: package 'foreach' was built under R version 4.4.3

Attaching package: 'foreach'

The following objects are masked from 'package:purrr':

```
accumulate, when
```

Loading required package: iterators

Warning: package 'iterators' was built under R version 4.4.3

Loading required package: parallel

```
library(RColorBrewer) # 用于自定义颜色  
# 加载集成模型包  
library(xgboost)
```

Warning: package 'xgboost' was built under R version 4.4.3

```
library(lightgbm)
```

Warning: package 'lightgbm' was built under R version 4.4.3

```
library(reticulate)
```

Warning: package 'reticulate' was built under R version 4.4.3

```
# loading  
df <- read.csv("C:/Users/ouyan/OneDrive/文档/BIOSTAT625/The_Cancer_data_1500_V2.csv")  
summary(df)
```

Age	Gender	BMI	Smoking
Min. :20.00	Min. :0.0000	Min. :15.00	Min. :0.0000
1st Qu.:35.00	1st Qu.:0.0000	1st Qu.:21.48	1st Qu.:0.0000
Median :51.00	Median :0.0000	Median :27.60	Median :0.0000
Mean :50.32	Mean :0.4907	Mean :27.51	Mean :0.2693
3rd Qu.:66.00	3rd Qu.:1.0000	3rd Qu.:33.85	3rd Qu.:1.0000

```

Max.    :80.00   Max.    :1.0000   Max.    :39.96   Max.    :1.0000
GeneticRisk      PhysicalActivity   AlcoholIntake      CancerHistory
Min.    :0.0000   Min.    :0.00241   Min.    :0.001215  Min.    :0.000
1st Qu.:0.0000   1st Qu.:2.43461   1st Qu.:1.210598  1st Qu.:0.000
Median  :0.0000   Median  :4.83432   Median  :2.382971  Median  :0.000
Mean    :0.5087   Mean    :4.89793   Mean    :2.417987  Mean    :0.144
3rd Qu.:1.0000   3rd Qu.:7.40990   3rd Qu.:3.585624  3rd Qu.:0.000
Max.    :2.0000   Max.    :9.99461   Max.    :4.987115  Max.    :1.000
Diagnosis
Min.    :0.0000
1st Qu.:0.0000
Median  :0.0000
Mean    :0.3713
3rd Qu.:1.0000
Max.    :1.0000

```

```
print(head(df))
```

	Age	Gender	BMI	Smoking	GeneticRisk	PhysicalActivity	AlcoholIntake
1	58	1	16.08531	0	1	8.146251	4.148219
2	71	0	30.82878	0	1	9.361630	3.519683
3	48	1	38.78508	0	2	5.135179	4.728368
4	34	0	30.04030	0	0	9.502792	2.044636
5	62	1	35.47972	0	0	5.356890	3.309849
6	27	0	37.10516	0	1	3.941905	2.324274
	CancerHistory	Diagnosis					
1		1					
2		0					
3		0		1			
4		0		0			
5		0		1			
6		0		0			

```
colSums(is.na(df))
```

Age	Gender	BMI	Smoking
0	0	0	0
GeneticRisk	PhysicalActivity	AlcoholIntake	CancerHistory
0	0	0	0
Diagnosis			
0			

```

library(tidyverse)
df <- df %>%
  mutate(
    Gender = factor(as.character(Gender),
                    levels = c("0", "1"),
                    labels = c("Male", "Female")),

```

```

Smoking = factor(as.character(Smoking),
                 levels = c("0", "1"),
                 labels = c("No","Yes")),

CancerHistory = factor(as.character(CancerHistory),
                      levels = c("0", "1"),
                      labels = c("No","Yes")),

GeneticRisk = factor(as.character(GeneticRisk),
                     levels = c("0","1","2"),
                     labels = c("Low","Medium","High")),

Diagnosis = factor(as.character(Diagnosis),
                   levels = c("0","1"),
                   labels = c("NoCancer","Cancer"))

)

print(head(df))

```

	Age	Gender	BMI	Smoking	GeneticRisk	PhysicalActivity	AlcoholIntake
1	58	Female	16.08531	No	Medium	8.146251	4.148219
2	71	Male	30.82878	No	Medium	9.361630	3.519683
3	48	Female	38.78508	No	High	5.135179	4.728368
4	34	Male	30.04030	No	Low	9.502792	2.044636
5	62	Female	35.47972	No	Low	5.356890	3.309849
6	27	Male	37.10516	No	Medium	3.941905	2.324274
				CancerHistory	Diagnosis		
1				Yes	Cancer		
2				No	NoCancer		
3				No	Cancer		
4				No	NoCancer		
5				No	Cancer		
6				No	NoCancer		

```

continuous_vars <- c("Age","BMI","PhysicalActivity","AlcoholIntake")
binary_vars <- c("Gender","Smoking","CancerHistory")
categorical_vars <- c("GeneticRisk")
target_var <- "Diagnosis"

```

```

# 定义用于划分和缩放的变量（确保这些变量已在你前面的代码中定义）
# continuous_vars <- c("Age","BMI","PhysicalActivity","AlcoholIntake")
# target_var <- "Diagnosis" # 假设这是你的目标变量

# 从你已经进行了因子转换的原始 df 开始
df_for_modeling <- df # 假设 'df' 是完成了因子转换后的数据框

# --- 1. 分割 x 和 y ---
X <- df_for_modeling %>% select(-all_of(target_var))
y <- df_for_modeling[[target_var]]

```

```

# --- 2. 划分训练集和测试集 (80/20 stratified split) ---
set.seed(42)
# 使用目标变量 y 进行分层抽样
train_index <- createDataPartition(y, p = 0.8, list = FALSE)

X_train <- X[train_index, ]
X_test <- X[-train_index, ]
y_train <- y[train_index]
y_test <- y[-train_index]

# --- 3. 标准化 (仅在训练集上学习参数) ---
# 构建预处理步骤：仅对训练集的连续特征进行标准化 (fit)
preprocessor <- preProcess(X_train[, continuous_vars], method = c("center", "scale"))

# 应用预处理：将训练集和测试集进行转换 (transform)
X_train_processed <- predict(preprocessor, X_train)
X_test_processed <- predict(preprocessor, X_test)

# --- 4. 创建最终的训练/测试数据集 (合并 x 和 y) ---
# 注意：X_train_processed 已经包含了所有因子变量
train_data <- cbind(Diagnosis = y_train, X_train_processed)
test_data <- cbind(Diagnosis = y_test, X_test_processed)

# 最终检查：确保目标变量是两级因子（前面已完成，但这一步保证了数据完整性）
train_data$Diagnosis <- factor(train_data$Diagnosis, levels = c("NoCancer", "Cancer"))
test_data$Diagnosis <- factor(test_data$Diagnosis, levels = c("NoCancer", "Cancer"))

# 现在 train_data 和 test_data 已经准备好用于模型训练
cat("数据划分和标准化已完成，准备进入模型训练阶段。\\n")

```

数据划分和标准化已完成，准备进入模型训练阶段。

```

# 1. 定义交叉验证参数 (5-Fold Stratified CV)
control <- trainControl(
  method = "cv",
  number = 5,
  classProbs = TRUE,          # Accuracy 不依赖概率，可保留
  summaryFunction = defaultSummary, # 改为默认 summary，这样会输出 Accuracy
  savePredictions = "final",
  verboseIter = FALSE
)

# 2. 定义模型列表 (使用caret支持的模型名称)
# 映射：Python名称 -> R caret 名称
# Logistic Regression -> glm (广义线性模型)
# Decision Tree -> rpart

```

```

# Random Forest -> rf
# Gradient Boosting -> gbm
# SVM -> svmRadial (径向核SVM)
# k-NN -> knn
# Naive Bayes -> nb
# LightGBM -> gbm (需要特殊设置或使用专用包)
# CatBoost -> catboost (需要使用专用包)

models_to_run <- c("glm", "rpart", "rf", "gbm", "svmRadial", "knn", "nb")
results_cv <- list()

# 3. 运行交叉验证
# 注册并行后端 (可选)
if (requireNamespace("doParallel", quietly = TRUE)) {
  cl <- makePSOCKcluster(detectCores() - 1)
  registerDoParallel(cl)
}

# 循环训练模型
for (model_name in models_to_run) {
  cat(paste("Training model:", model_name, "...\\n"))

  # Rpart (决策树) 的调参网格
  if (model_name == "rpart") {
    tune_grid <- expand.grid(cp = seq(0.01, 0.1, 0.01))
  } else {
    tune_grid <- NULL # 使用默认调参网格
  }

  # 训练模型
  model_fit <- train(
    Diagnosis ~.,
    data = train_data,
    method = model_name,
    trControl = control,
    metric = "Accuracy", # 使用准确率作为选择指标
    tuneGrid = tune_grid
  )

  results_cv[[model_name]] <- model_fit

  # 打印结果
  print(model_fit$results[order(model_fit$results$Accuracy, decreasing = TRUE), ][1, ])
}

```

```

Training model: glm ...
parameter Accuracy      Kappa AccuracySD      KappaSD
1      none 0.8792531 0.738078 0.02133139 0.04664393
Training model: rpart ...
cp  Accuracy      Kappa AccuracySD      KappaSD

```

```

1 0.01 0.8509647 0.6745558 0.02448481 0.0489391
Training model: rf ...
  mtry Accuracy      Kappa AccuracySD      KappaSD
2      5 0.9084094 0.8015616 0.02716388 0.05779453

Training model: gbm ...
  Iter   TrainDeviance   ValidDeviance   StepSize   Improve
    1       1.2682           nan     0.1000   0.0266
    2       1.2253           nan     0.1000   0.0220
    3       1.1884           nan     0.1000   0.0164
    4       1.1617           nan     0.1000   0.0135
    5       1.1385           nan     0.1000   0.0102
    6       1.1173           nan     0.1000   0.0086
    7       1.0931           nan     0.1000   0.0114
    8       1.0763           nan     0.1000   0.0059
    9       1.0533           nan     0.1000   0.0104
   10       1.0343           nan     0.1000   0.0093
   20       0.8896           nan     0.1000   0.0053
   40       0.7116           nan     0.1000   0.0024
   60       0.6084           nan     0.1000   0.0015
   80       0.5477           nan     0.1000   0.0007
  100       0.5031           nan     0.1000   0.0008

  shrinkage interaction.depth n.minobsinnode n.trees   Accuracy      Kappa
5          0.1                  2          10        100 0.9492185 0.8901224
  AccuracySD      KappaSD
5 0.01537361 0.03309005

Training model: svmRadial ...
  sigma C Accuracy      Kappa AccuracySD      KappaSD
3 0.08401214 1 0.870121 0.7180163 0.01173217 0.02381047

Training model: knn ...
  k Accuracy      Kappa AccuracySD      KappaSD
2 7 0.8110028 0.5645195 0.01555807 0.03740757

Training model: nb ...
  usekernel fL adjust   Accuracy      Kappa AccuracySD      KappaSD
1 FALSE 0          1 0.7935201 0.5255632 0.01201641 0.02750503

```

```

# 停止并行后端 (如果已注册)
if (exists("cl")) {
  stopCluster(cl)
  registerDoSEQ()
}

# 4. 汇总和选择最佳模型
results_summary <- data.frame(
  Model = character(),
  CV_Accuracy_Mean = numeric(),
  CV_Accuracy_Std = numeric(), # 修正此处的列名
  stringsAsFactors = FALSE
)

```

```

for (name in names(results_cv)) {
  best_tune <- results_cv[[name]]$results[order(results_cv[[name]]$results$Accuracy, decreasing
  results_summary <- results_summary %>%
    add_row(
      Model = name,
      CV_Accuracy_Mean = best_tune$Accuracy,
      CV_Accuracy_Std = best_tune$AccuracySD #关键修正：使用 AccuracySD
    )
}

# 3. 打印 CV 汇总
print("== Cross-Validation Results ==")

```

[1] "== Cross-Validation Results =="

```
print(results_summary)
```

	Model	CV_Accuracy_Mean	CV_Accuracy_Std
1	glm	0.8792531	0.02133139
2	rpart	0.8509647	0.02448481
3	rf	0.9084094	0.02716388
4	gbm	0.9492185	0.01537361
5	svmRadial	0.8701210	0.01173217
6	knn	0.8110028	0.01555807
7	nb	0.7935201	0.01201641

# 4. 选择最佳模型 (Accuracy 最大)

```

best_model_name <- results_summary$Model[which.max(results_summary$CV_Accuracy_Mean)]
best_model <- results_cv[[best_model_name]]
cat(paste("\nBest model:", best_model_name, "with mean CV Accuracy=",
         max(results_summary$CV_Accuracy_Mean), "\n"))

```

Best model: gbm with mean CV Accuracy= 0.949218533886584

```

# 1. 测试集评估
test_results_list <- list()

for (model_name in names(results_cv)) {
  model_fit <- results_cv[[model_name]]

  # 预测
  y_pred <- predict(model_fit, newdata = test_data)

  # 计算混淆矩阵
  cm <- confusionMatrix(y_pred, test_data$Diagnosis, positive = "Cancer")

  # 提取指标
}
```

```

acc <- cm$overall['Accuracy']
prec <- cm$byClass['Pos Pred Value'] # Precision
rec <- cm$byClass['Recall']           # Recall
f1 <- cm$byClass['F1']

test_results_list[[model_name]] <- data.frame(
  Model = model_name,
  Test_Accuracy = acc,
  Test_Precision = prec,
  Test_Recall = rec,
  Test_F1 = f1
)

# 打印混淆矩阵 (对应Python中的 ConfusionMatrixDisplay plot)
if (model_name == best_model_name) {
  cat(paste("\n==== Confusion Matrix -", model_name, "====\n"))
  print(cm$table)

  # 混淆矩阵可视化
  cm_df <- as.data.frame(cm$table)

  # 使用 ggplot2 绘制热图形式的混淆矩阵
  ggplot(cm_df, aes(x = Prediction, y = Reference, fill = Freq)) +
    geom_tile(color = "white") +
    scale_fill_gradient(low = "white", high = "#336699") +
    geom_text(aes(label = Freq), vjust = 1, fontface = "bold", size = 5) +
    labs(title = paste("Confusion Matrix -", model_name),
         x = "Predicted Label", y = "True Label") +
    theme_minimal() +
    theme(plot.title = element_text(face = "bold", size = 14), legend.position = "none")
}
}

```

```

==== Confusion Matrix - gbm ====
      Reference
Prediction NoCancer Cancer
NoCancer      185     11
Cancer        3     100

```

```

# 2. 汇总测试结果
test_results_df <- do.call(rbind, test_results_list)

print("\n==== Test Set Results ===")

```

```
[1] "\n==== Test Set Results ==="
```

```
print(test_results_df)
```

	Model	Test_Accuracy	Test_Precision	Test_Recall	Test_F1
glm	glm	0.8829431	0.8584906	0.8198198	0.8387097
rpart	rpart	0.8762542	0.8557692	0.8018018	0.8279070
rf	rf	0.9230769	0.9150943	0.8738739	0.8940092
gbm	gbm	0.9531773	0.9708738	0.9009009	0.9345794
svmRadial	svmRadial	0.8829431	0.8800000	0.7927928	0.8341232
knn	knn	0.8193980	0.9130435	0.5675676	0.7000000
nb	nb	0.7892977	0.8428571	0.5315315	0.6519337

```
# 3. 保存结果
#write.csv(results_summary, 'cross_validation_results_R.csv', row.names = FALSE)
#write.csv(test_results_df, 'test_set_results_R.csv', row.names = FALSE)
#cat("\nResults saved to 'cross_validation_results_R.csv' and 'test_set_results_R.csv'.\n")
```

```
library(dplyr)

df %>%
  group_by(Diagnosis) %>%
  summarise(
    Count = n(),
    Percentage = n() / nrow(df) * 100
  ) %>%
  ungroup() %>%
  mutate(
    # 格式化百分比为两位小数
    Percentage_Formatted = paste0(round(Percentage, 2), "%")
  )
```

```
# A tibble: 2 × 4
Diagnosis Count Percentage Percentage_Formatted
<fct>     <int>      <dbl> <chr>
1 NoCancer   943       62.9 62.87%
2 Cancer     557       37.1 37.13%
```

```
set.seed(42)

K <- 5
folds <- createFolds(train_data$Diagnosis, k = K, list = TRUE)

# 用来存放 OOF stacking 特征
oof_rf_prob <- rep(NA, nrow(train_data))
oof_gbm_prob <- rep(NA, nrow(train_data))

# 5-fold stacking
for (i in 1:K) {
  cat("Processing fold", i, "\n")

  idx_valid <- folds[[i]]
  idx_train <- setdiff(seq_len(nrow(train_data)), idx_valid)
```

```

train_fold <- train_data[idx_train, ]
valid_fold <- train_data[idx_valid, ]

# RF
rf_model <- train(
  Diagnosis ~ ., data = train_fold, method = "rf",
  trControl = trainControl(method = "cv", number = 3),
  tuneLength = 5
)

# GBM
gbm_model <- train(
  Diagnosis ~ ., data = train_fold, method = "gbm",
  trControl = trainControl(method = "cv", number = 3),
  verbose = FALSE
)

# OOF predictions
oof_rf_prob[idx_valid] <- predict(rf_model, valid_fold, type = "prob")[, "Cancer"]
oof_gbm_prob[idx_valid] <- predict(gbm_model, valid_fold, type = "prob")[, "Cancer"]
}

```

Processing fold 1  
 Processing fold 2  
 Processing fold 3  
 Processing fold 4  
 Processing fold 5

```

# 生成无泄漏 stacked 特征
stacked_train <- data.frame(
  rf_prob = oof_rf_prob,
  gbm_prob = oof_gbm_prob,
  y = as.numeric(train_data$Diagnosis == "Cancer")
)

```

```
library(glmnet)
```

Warning: package 'glmnet' was built under R version 4.4.3

Loading required package: Matrix

Warning: package 'Matrix' was built under R version 4.4.3

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyঃ':

```
expand, pack, unpack
```

```
Loaded glmnet 4.1-10
```

```
X_stack <- as.matrix(stacked_train[, c("rf_prob", "gbm_prob")])
y_stack <- stacked_train$y

stack_lr <- cv.glmnet(
  X_stack, y_stack,
  family = "binomial",
  alpha = 1,           # LASSO: prevent over fitting
  nfolds = 5
)

cat("Selected lambda:", stack_lr$lambda.min, "\n")
```

```
Selected lambda: 0.003861134
```

```
coef(stack_lr, s = "lambda.min")
```

```
3 x 1 sparse Matrix of class "dgCMatrix"
  lambda.min
(Intercept) -5.1876271
rf_prob      0.6420093
gbm_prob     10.9185016
```

```
rf_full <- train(
  Diagnosis ~ ., data = train_data, method = "rf",
  trControl = trainControl(method = "none"),
  tuneGrid = data.frame(mtry = floor(sqrt(ncol(train_data) - 1))),
  ntree = 500
)

gbm_full <- train(
  Diagnosis ~ ., data = train_data, method = "gbm",
  trControl = trainControl(method = "none"),
  verbose = FALSE,
  tuneGrid = data.frame(
    n.trees = 200,
    interaction.depth = 3,
    shrinkage = 0.05,
    n.minobsinnode = 10
  )
)
```

```
rf_test_prob <- predict(rf_full, newdata = test_data, type = "prob")[, "Cancer"]
gbm_test_prob <- predict(gbm_full, newdata = test_data, type = "prob")[, "Cancer"]

stacked_test <- as.matrix(data.frame(
```

```

rf_prob = rf_test_prob,
gbm_prob = gbm_test_prob
))

final_probs <- predict(stack_lr, newx = stacked_test, s = "lambda.min", type = "response")
final_pred <- ifelse(final_probs > 0.5, "Cancer", "NoCancer")
final_pred <- factor(final_pred, levels = c("NoCancer", "Cancer"))

# Confusion Matrix
confusionMatrix(final_pred, test_data$Diagnosis)

```

### Confusion Matrix and Statistics

		Reference	
		NoCancer	Cancer
NoCancer	181	10	
Cancer	7	101	

Accuracy : 0.9431  
95% CI : (0.9105, 0.9665)  
No Information Rate : 0.6288  
P-Value [Acc > NIR] : <2e-16

Kappa : 0.8775

McNemar's Test P-Value : 0.6276

Sensitivity : 0.9628  
Specificity : 0.9099  
Pos Pred Value : 0.9476  
Neg Pred Value : 0.9352  
Prevalence : 0.6288  
Detection Rate : 0.6054  
Detection Prevalence : 0.6388  
Balanced Accuracy : 0.9363

'Positive' Class : NoCancer

```

train_pred <- predict(stack_lr, newx = X_stack, s = "lambda.min", type = "response")
# Convert matrix into vector
train_pred_vector <- as.numeric(train_pred[, 1]) # 提取第一列并转为 numeric 向量
final_probs_vector <- as.numeric(final_probs[, 1])

# ROC curve
roc_curve <- roc(test_data$Diagnosis, final_probs_vector)

```

Setting levels: control = NoCancer, case = Cancer

```
Setting direction: controls < cases
```

```
auc(roc_curve)
```

```
Area under the curve: 0.9492
```

```
library(MLmetrics)
```

```
Warning: package 'MLmetrics' was built under R version 4.4.3
```

```
Attaching package: 'MLmetrics'
```

```
The following objects are masked from 'package:caret':
```

```
MAE, RMSE
```

```
The following object is masked from 'package:base':
```

```
Recall
```

```
# 将 factor 转成 numeric (1 = Cancer, 0 = NoCancer)
y_true <- ifelse(test_data$Diagnosis == "Cancer", 1, 0)
y_pred <- ifelse(final_pred == "Cancer", 1, 0)

Precision <- Precision(y_pred, y_true)
Recall    <- Recall(y_pred, y_true)
F1        <- F1_Score(y_pred, y_true)

cat("Precision:", Precision, "\n")
```

```
Precision: 0.962766
```

```
cat("Recall:", Recall, "\n")
```

```
Recall: 0.947644
```

```
cat("F1 Score:", F1, "\n")
```

```
F1 Score: 0.9551451
```

```
# (i) Evaluate over fitting
train_pred <- predict(stack_lr, newx = X_stack, s = "lambda.min", type = "response")

auc_train <- auc(roc(train_data$Diagnosis, train_pred_vector))
```

```
Setting levels: control = NoCancer, case = Cancer
```

```
Setting direction: controls < cases
```

```
auc_test <- auc(roc(test_data$Diagnosis, final_probs_vector))
```

```
Setting levels: control = NoCancer, case = Cancer
```

```
Setting direction: controls < cases
```

```
cat("Train AUC:", auc_train, "\n")
```

```
Train AUC: 0.961824
```

```
cat("Test AUC:", auc_test, "\n")
```

```
Test AUC: 0.9492045
```

```
cat("Gap =", auc_train - auc_test, "\n")
```

```
Gap = 0.01261949
```

```
# 1. 确保 Stacked Train Vector 已转换
# train_pred <- predict(stack_lr, newx = X_stack, s = "lambda.min", type = "response")
# train_pred_vector <- as.numeric(train_pred[, 1])

# 2. 获取 GLM 在测试集上的概率预测
glm_model <- results_cv[["glm"]]
glm_test_probs <- predict(glm_model, newdata = test_data, type = "prob")[, "Cancer"]

# 3. 计算 GLM 的 ROC 曲线
roc_glm <- roc(test_data$Diagnosis, glm_test_probs)
```

```
Setting levels: control = NoCancer, case = Cancer
```

```
Setting direction: controls < cases
```

```
# --- 绘图 ---
# 绘制 Stacked Train ROC (作为基础图)
plot(roc(train_data$Diagnosis, train_pred_vector),
      col = "blue",
      main = "Stacked vs. Logistic Regression ROC Comparison",
      # 设置 y 轴和 x 轴标签
      xlab = "False Positive Rate (1 - Specificity)",
      ylab = "True Positive Rate (Sensitivity)")
```

```
Setting levels: control = NoCancer, case = Cancer
```

```
Setting direction: controls < cases
```

```
# 叠加 Stacked Test ROC
plot(roc(test_data$Diagnosis, final_probs_vector),
  col = "red",
  add = TRUE)
```

Setting levels: control = NoCancer, case = Cancer  
Setting direction: controls < cases

```
# 叠加 GLM Test ROC
plot(roc_glm,
  col = "purple",
  add = TRUE)

# 4. 添加图例和随机猜测线
abline(a = 0, b = 1, lty = 2, col = "gray") # 随机猜测线


legend("bottomright",
  legend = c(
    paste("Stacked Train AUC:", round(auc(roc(train_data$Diagnosis, train_pred_vector)), 4)),
    paste("Stacked Test AUC:", round(auc(roc(test_data$Diagnosis, final_probs_vector)), 4)),
    paste("Logistic Regression Test AUC:", round(auc(roc_glm), 4))
  ),
  col = c("blue", "red", "purple"),
  lwd = 2)
```

Setting levels: control = NoCancer, case = Cancer  
Setting direction: controls < cases

Setting levels: control = NoCancer, case = Cancer

Setting direction: controls < cases

## Stacked vs. Logistic Regression ROC Comparison



```
cat("\nLogistic Regression Test AUC:", auc(roc_glm), "\n")
```

Logistic Regression Test AUC: 0.9261549

```
# (ii) Evaluate the stratified split  
prop.table(table(train_data$Diagnosis))
```

```
NoCancer      Cancer  
0.6286428  0.3713572
```

```
prop.table(table(test_data$Diagnosis))
```

```
NoCancer      Cancer  
0.6287625  0.3712375
```


```
library(ggplot2)  
  
ggplot(train_data, aes(x = Diagnosis)) +  
  geom_bar(aes(y =(..count..)/sum(..count..)), fill = "steelblue") +
```

```

geom_text(aes(y =(..count..)/sum(..count..), label = scales::percent(..count..)/sum(..count..)
               stat = "count", vjust = -0.5) +
ylab("Proportion") +
ggtitle("Training Set Class Distribution") +
theme_minimal()

```

Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2 3.4.0.  
 i Please use `after\_stat(count)` instead.




```

ggplot(test_data, aes(x = Diagnosis)) +
  geom_bar(aes(y =(..count..)/sum(..count..)), fill = "steelblue") +
  geom_text(aes(y =(..count..)/sum(..count..), label = scales::percent(..count..)/sum(..count..)
                stat = "count", vjust = -0.5) +
ylab("Proportion") +
ggtitle("Testing Set Class Distribution") +
theme_minimal()

```

## Testing Set Class Distribution



```
# 3. (iii) Calibration Plot
# 创建数据框
df_calib <- data.frame(
  prob = final_probs_vector,
  true = y_true
)

# 2. 将预测概率分组
n_bins <- 6
df_calib$bin <- cut(df_calib$prob, breaks = seq(0,1,length.out = n_bins+1), include.lowest = TRUE)


calib_summary <- df_calib %>%
  group_by(bin) %>%
  summarise(
    mean_pred = mean(prob),
    obs_rate = mean(true),
    n         = n()
  )

# 3. 绘制校准曲线
ggplot(calib_summary, aes(x = mean_pred, y = obs_rate)) +
  geom_point(aes(size = n), color = "blue") +          # 每个 bin 点大小按样本数
  geom_line(color = "blue") +
```

```

# 3. 添加文本标签：显示每个点的样本量 n
# vjust/hjust 用于微调文本位置，使其不与点重叠
geom_text(aes(label = n), vjust = -1.5, hjust = 0.5, size = 3.5, color = "black") +
geom_abline(slope = 1, intercept = 0, linetype = "dashed", color = "red") +
xlim(0, 1) + ylim(0, 1) +
xlab("Mean Predicted Probability") +
ylab("Observed Event Rate") +
ggtitle("Stacked Logistic Regression Calibration Plot") +
theme_minimal() +
scale_size_continuous(name = "Samples per bin")

```



```

# 4. Calculate Brier Score
brier_score <- mean((final_probs_vector - y_true)^2)
cat("Brier Score:", round(brier_score, 4), "\n")

```

Brier Score: 0.0525

```

# (iv) Evaluate feature importance / Explainability (模型可解释性)
# rf 特征重要性
rf_imp <- varImp(rf_full)$importance
rf_imp %>% arrange(desc(Overall))

```

	Overall
CancerHistoryYes	100.00000
BMI	83.93317
AlcoholIntake	80.48403
GeneticRiskHigh	79.50337
PhysicalActivity	75.92005
Age	74.19066
GenderFemale	42.45197
SmokingYes	26.34592
GeneticRiskMedium	0.00000

```
# GBM 特征重要性
gbm_imp <- summary(gbm_full$finalModel, n.trees = gbm_full$bestTune$n.trees, plotit = FALSE)
print(gbm_imp)
```

	var	rel.inf
CancerHistoryYes	CancerHistoryYes	20.90859506
GeneticRiskHigh	GeneticRiskHigh	17.85936796
AlcoholIntake	AlcoholIntake	11.10167064
GenderFemale	GenderFemale	10.90965907
BMI	BMI	10.68898422
PhysicalActivity	PhysicalActivity	10.60143407
Age	Age	10.51628965
SmokingYes	SmokingYes	7.33820859
GeneticRiskMedium	GeneticRiskMedium	0.07579074

```
coef(stack_lr, s = "lambda.min") # 查看各 base model 概率的权重
```

```
3 x 1 sparse Matrix of class "dgCMatrix"
      lambda.min
(Intercept) -5.1876271
rf_prob       0.6420093
gbm_prob      10.9185016
```

```
library(dplyr)
library(tidyr)

rf_data <- data.frame(
  Feature = c("CancerHistoryYes", "BMI", "AlcoholIntake", "GeneticRiskHigh", "PhysicalActivity",
             "Age", "GenderFemale", "SmokingYes", "GeneticRiskMedium"),
  RF_Importance = c(100.00, 83.93, 80.48, 79.50, 75.92, 74.19, 42.45, 26.35, 0.00)
)

# --- 模拟您提供的 GBM 数据 ---
gbm_data <- data.frame(
  Feature = c("CancerHistoryYes", "GeneticRiskHigh", "AlcoholIntake", "GenderFemale", "BMI",
             "PhysicalActivity", "Age", "SmokingYes", "GeneticRiskMedium"),
  GBM_Importance = c(20.91, 17.86, 11.10, 10.91, 10.69, 10.60, 10.52, 7.34, 0.08)
)
```

```

# --- 统一格式和合并 ---
combined_imp <- full_join(rf_data, gbm_data, by = "Feature") %>%
  # 简化/统一特征名称
  mutate(
    Feature_Clean = case_when(
      grepl("CancerHistory", Feature) ~ "Cancer History (Yes)",
      grepl("GeneticRiskHigh", Feature) ~ "Genetic Risk (High)",
      grepl("GenderFemale", Feature) ~ "Gender (Female)",
      grepl("SmokingYes", Feature) ~ "Smoking (Yes)",
      grepl("GeneticRiskMedium", Feature) ~ "Genetic Risk (Medium)",
      TRUE ~ Feature # 保留其他变量原名
    )
  ) %>%
  select(Feature_Clean, RF_Importance, GBM_Importance) %>%
  arrange(desc(RF_Importance)) # 以 RF 重要性作为主排序

# 打印最终结果
print(combined_imp)

```

	Feature_Clean	RF_Importance	GBM_Importance
1	Cancer History (Yes)	100.00	20.91
2	BMI	83.93	10.69
3	AlcoholIntake	80.48	11.10
4	Genetic Risk (High)	79.50	17.86
5	PhysicalActivity	75.92	10.60
6	Age	74.19	10.52
7	Gender (Female)	42.45	10.91
8	Smoking (Yes)	26.35	7.34
9	Genetic Risk (Medium)	0.00	0.08

```

# 定义一个函数，用于计算 Stacked 模型的最终概率
stacked_predict_prob <- function(object, newdata) {
  # 1. Base Learner Predictions
  rf_test_prob <- predict(rf_full, newdata = newdata, type = "prob")[, "Cancer"]
  gbm_test_prob <- predict(gbm_full, newdata = newdata, type = "prob")[, "Cancer"]
  # 🚨 Sanity Check 1: 检查 Base Learner 预测长度是否匹配输入
  if (length(rf_test_prob) != nrow(newdata) || length(gbm_test_prob) != nrow(newdata)) {
    warning("Sanity Check Failed: Base Learner prediction length does not match input data rows. Check for NA/missing rows due to new factor levels.")
  }
  # 理想情况下，应该停止执行
  # stop("Inconsistent prediction length detected.")
}

# 2. Create Stacked Feature Matrix (Level 1)
if (any(is.na(rf_test_prob)) || any(is.na(gbm_test_prob))) {
  warning("Sanity Check Failed: Base Learner probabilities contain NA values.")
}

stacked_test <- as.matrix(data.frame(
  rf_prob = rf_test_prob,
  gbm_prob = gbm_test_prob
))

```

```

))
# 3. Final Prediction using Meta Learner
final_probs <- predict(stack_lr, newx = stacked_test, s = "lambda.min", type = "response")
# ! Sanity Check 2: 检查最终概率是否有效
if (any(final_probs < 0) || any(final_probs > 1) || any(is.na(final_probs))) {
  warning("Sanity Check Failed: Final probabilities contain values outside [0, 1] or contain NA.")
  # stop("Invalid probability detected.")
}
# 返回概率向量
return(as.numeric(final_probs))
}

```

```
library(iml)
```

Warning: package 'iml' was built under R version 4.4.3

```

# --- 1. 创建 Stacked 模型的解释器对象 ---
# 传入一个占位符模型 (rf_full) 作为 model, 但关键是传递自定义预测函数
predictor_stacked <- Predictor$new(
  # model = rf_full, # 占位符
  data = train_data[,-which(names(train_data)=="Diagnosis")],
  y = train_data$Diagnosis,
  predict.function = stacked_predict_prob, # 关键: 指定 Stacked 预测函数
  class = "Cancer" # 明确指定要解释的类别
)


# --- 2. 选取要解释的样本 (测试集第 50 个样本) ---
x_interest_test_sample <- test_data[50, -which(names(test_data)=="Diagnosis")]

# --- 3. 计算 Stacked 模型的 SHAP 值 ---
shapley_stacked <- Shapley$new(
  predictor = predictor_stacked,
  x.interest = x_interest_test_sample
)

# --- 4. 绘图 ---
shapley_stacked$plot()

```

Actual prediction: 1.00  
Average prediction: 0.38



```
# Sanity check
# 1. 创建一个测试数据框 (例如, 使用训练集的前 10 行)
test_data_sample <- train_data[1:10, -which(names(train_data) == "Diagnosis")]

# 2. 运行你的 stacked 预测函数
test_probs <- stacked_predict_prob(object = rf_full, newdata = test_data_sample)

# 3. 严格检查最大值
max_prob <- max(test_probs, na.rm = TRUE)
min_prob <- min(test_probs, na.rm = TRUE)

print(paste("Maximum calculated probability:", max_prob))
```

```
[1] "Maximum calculated probability: 0.997020085345288"
```

```
print(paste("Minimum calculated probability:", min_prob))
```

```
[1] "Minimum calculated probability: 0.00626285133180356"
```

```
library(iml)
library(ggplot2)
# Validated that our stacked logistic regression has learned non-linear relationship
```

```


# --- 1. 使用 iml::FeatureEffect 计算 PDP 数据 ---
# method = "pdp" 是偏依赖图
effect_bmi <- FeatureEffect$new(
  predictor = predictor_stacked,
  feature = "BMI",
  method = "pdp"
)

# --- 2. 绘制 PDP 图 ---
# iml 的 plot() 方法会自动处理 PDP 的标签和尺度
# 确保 y 轴在 0 到 1 之间
plot_pdp_bmi <- effect_bmi$plot() +
  labs(title = "Stacked Model PDP (BMI)",
       y = "Predicted Cancer Probability (0-1)") +
  theme_minimal()

print(plot_pdp_bmi)

```

Stacked Model PDP (BMI)



```

# AlcoholIntake plot
effect_alch <- FeatureEffect$new(
  predictor = predictor_stacked,
  feature = "AlcoholIntake",
  method = "pdp"
)


```

```

)
plot_pdp_alch <- effect_alch$plot() +
  labs(title = "Stacked Model PDP (AlcoholIntake)",
       y = "Predicted Cancer Probability (0-1)") +
  theme_minimal()

print(plot_pdp_alch)

```



```

# (v) Deviance and QQ-plot as diagnosis
# --- 1. 准备训练集预测概率 ---
X_train_stack <- as.matrix(data.frame(
  rf_prob = predict(rf_full, newdata = train_data, type = "prob")[, "Cancer"],
  gbm_prob = predict(gbm_full, newdata = train_data, type = "prob")[, "Cancer"]
))
y_train <- ifelse(train_data$Diagnosis == "Cancer", 1, 0)

train_pred_prob <- as.vector(predict(stack_lr, newx = X_train_stack, s = "lambda.min", type = "re")
# --- 2. 计算残差 ---
# Deviance residuals
dev_resid <- ifelse(y_train == 1,
                      sqrt(-2 * log(train_pred_prob)),
                      -sqrt(-2 * log(1 - train_pred_prob)))

```

```
# Standardized residuals
std_resid <- scale(dev_resid)

# --- 2. 偏差残差 vs 拟合值图 (Residuals vs Fitted) ---
# 检查残差是否随机分布在 0 附近
residual_data <- data.frame(Fitted = train_pred_prob, DevianceResid = dev_resid)


plot_resid_fitted <- ggplot(residual_data, aes(x = Fitted, y = DevianceResid)) +
  geom_point(alpha = 0.5, color = "#1E88E5") + # 蓝色散点
  geom_smooth(method = "loess", color = "darkorange", se = FALSE, linetype = "solid") + # 添加平滑
  geom_hline(yintercept = 0, color = "red", linetype = "dashed", size = 1) +
  labs(title = "Stacked Model: Deviance Residuals vs Fitted Probabilities",
       x = "Fitted Probabilities",
       y = "Deviance Residuals") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5, face = "bold"))
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.  
i Please use `linewidth` instead.

```
print(plot_resid_fitted)
```

```
`geom_smooth()` using formula = 'y ~ x'
```

## Stacked Model: Deviance Residuals vs Fitted Probabilities



```
# --- 3. 标准化残差 Q-Q Plot ---
# 检查偏差残差是否近似正态分布
qq_data <- data.frame(StdResid = std_resid)

plot_qq <- ggplot(qq_data, aes(sample = StdResid)) +
  stat_qq(alpha = 0.6, color = "#2E7D32") + # 绿色散点
  stat_qq_line(color = "red", linetype = "dashed", size = 1) +
  labs(title = "Stacked Model: Q-Q Plot of Standardized Deviance Residuals",
       x = "Theoretical Quantiles",
       y = "Sample Quantiles") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5, face = "bold"))

print(plot_qq)
```

### Stacked Model: Q-Q Plot of Standardized Deviance Residuals

